

ENGINEERING STATUS AND FORWARD PLAN

AI Decision Firewall - Phase 3 simulation-only operational MVP candidate

Prior evidence baseline	0.2.0-alpha.5 commit f7f6b5c Phase 2.4
Published predecessor	0.2.0-alpha.6 commit 854b15c CI and Dependency Graph succeeded
Phase 3 candidate	0.3.0-alpha.1 288/288 local regression exact commit and CI pending
Review date	2026-08-15
Authority	In-memory synthetic simulation only; no live connector or operational authority
Evidence boundary	Phase 3 local CE-1 only; P2-CE-005 remains CE-0 NOT_EVALUATED

Executive assessment

The program has two distinct current boundaries.

First, Phase 2.5 is published on `main` at exact Commit `854b15c56397a81de6326b719d3d7d1dc847608f`. GitHub CI and Dependency Graph checks passed for that exact commit. It retains the local 222/222 Phase 2.5 technical result, separate 9/9 public-site result, and 231/231 then-current repository aggregate. `P2-CE-005` was not executed or published and remains `CE-0 NOT_EVALUATED`; green package CI is not campaign evidence.

Second, Phase 3 is a working **local 0.3.0-alpha.1 simulation-only operational-MVP candidate**. It adds a raw external request, opaque-credential identity resolution, signed target-bound evidence, trusted policy/target context, deterministic four-way decision, consequence evaluation, exact-scope single-use authorization, mandatory in-memory broker, functionally separate same-project target readback, bounded human approval, lifecycle audit, metrics, two SOC demonstrations, and a 46-scenario adversarial corpus. The current checkout passed 57/57 focused Phase 3 tests; both demo acceptance checks reported PASS; the corpus reported 46/46; and the full repository regression passed 288/288. No Phase 3 commit, publication, or exact-commit CI result exists yet.

All Phase 3 observations are CE-1 implementation-conformance evidence over synthetic inputs and effects. They do not establish live containment, operational efficacy or calibration, production safety, statistical risk, or external independence.

Decision table

State	Current truth
Published Phase 2.5 code/package	Exact Commit 854b15c56397a81de6326b719d3d7d1dc847608f on <code>main</code> ; exact-commit CI and Dependency Graph green
Published Phase 2 evidence	P2-CE-001 through P2-CE-004, version bound to their records and limitations
Phase 2.5 campaign evidence	P2-CE-005 CE-0 NOT_EVALUATED; no execution, result ledger, pass rate, CE-2 record, or evidence-only Commit B
Phase 3 implementation	Working local candidate; raw request through optional synthetic effect and separate readback implemented
Phase 3 local T&E	Required high-risk/no-effect and low-risk/verified-effect demo acceptance PASS; deterministic corpus 46/46; focused tests 57/57
Phase 3 repository aggregate	288/288 passed in the settled local checkout; exact-commit CI pending
Phase 3 external gates	Final exact commit, explicit publication authorization, GitHub CI and Dependency Graph
Data/action boundary	Synthetic only; no historical organizational data, live feed, test tenant, production connector, credential, or live action

What Phase 3 adds




Figure 1. Phase 3 simulation-only operational MVP control flow and accountability boundary.

Request and trust plane

- Closed v0.3.0 raw-request and policy schemas with strict duplicate/non-finite, size/depth, version, field, and timestamp handling.

- Opaque invocation credential resolution to a signed `ResolvedPrincipal`, plus trusted policy registries for agent authority, source properties, action constraints, target criticality/dependencies/consequences, and firewall-owned time.
- Runtime HMAC evidence attestation over source, provenance, observation time, content, semantic support/contradiction, relevance, and subject target.
- Freshness-at-decision, reliability, relevance, corroboration, conflict, missing-source, content-integrity, and poisoned-text assessment.

Decision and consequence plane

- Exact `ALLOW`, `DENY`, `ESCALATE`, and `ALLOW_CONSTRAINED` outcomes.
- Structured stable reason codes, applicable rules, evidence/authority/consequence findings, constraints, policy digest, and decision-context digest.
- Explicit consequence evaluation for reversibility, criticality, dependencies, cascade, blast radius, downtime, mission, safety, availability, and required human authority.
- Recommendation/confidence separation: agent output remains advisory and never grants authority.
- Functionally separate deterministic decision verification before token issue.
- Code-owned policy safety invariants for unique closed rule precedence, evidence trust/strength/corroboration/zero-conflict floors, severe-consequence approval floors, and Tier-0 treatment for every domain controller.

Authorization, execution, and verification plane

- HMAC authorization binding issuer instance, request, decision, agent, action, target, canonical permitted parameters, issue/expiry time, policy identity and digest, decision context, target precondition, nonce, and signature.
- Atomic process-local single-use consumption, including failed attempts; sequential, concurrent, prior-instance, expired, altered, and wrong-scope reuse is rejected.
- Mandatory broker and in-memory target mutation guarded by a private execution capability and a target-state precondition.
- Separate read-only target observation; broker-reported success cannot produce `VERIFIED`. Final classes include `FAILED`, `PARTIAL`, `UNEXPECTED_EFFECT`, and `ROLLBACK_REQUIRED`.

- Signed registered expiring human-approval requirement. A separately resolved

opaque human credential can create one atomically recorded exact-scope signed reevaluation receipt; the approval gate cannot mint a token, cause reevaluation, or invoke the broker.

Audit and observability

- Correlated hash-linked request lifecycle through validation/decision,

authorization or suppression, broker/attempt, verification, and final state.

- Explicit nonexecuting and broker-rejected terminal paths without fabricated

downstream events.

- Exact executed-path request/decision/token/attempt/target/state/effect/

verification correlation. A post-effect prewrite audit failure closes through one honest POST_EFFECT_ACCOUNTING_FAILURE, returns ROLLBACK_REQUIRED, and reconciles the decision/verification-failure counters exactly once; it does not automate rollback.

- Decision/rule/conflict/authorization/broker/verification/latency metrics.

- Demo and corpus outputs omit reusable token signatures and runtime keys and

refuse to overwrite nonempty destinations.

Demonstrated mission paths

Demonstration	Trusted conditions	Observed local result
Tier-0 domain controller	AI requests isolation at confidence 0.96; stale and conflicting evidence; authentication-service dependency and cascading consequence; agent lacks Tier-0 authority	ESCALATE; bound human-review requirement; no authorization, broker call, or state change
Low-criticality workstation	Credential-resolved workstation-containment authority; fresh corroborated HMAC-attested evidence; reversible bounded action	ALLOW; one exact-scope authorization and synthetic broker attempt; functionally separate same-project observation VERIFIED

The deterministic corpus adds 46 declared canonical, evidence, authority, consequence, token/replay/bypass, broker/verifier, metamorphic, and combined attack cases. All 46 matched their project-controlled expectations locally. This is complete coverage of the declared corpus, not exhaustive security coverage or a statistical failure estimate.

Safety defects found and closed

Adversarial review found and closed release-blocking defects across cascading consequence and signed subject-target evidence binding; opaque credential and trust-domain key handling; polymorphic/late-mutable security values; machine policy safety floors; request/token/verifier/approval replay and failure atomicity; injected clock, identifier, verifier, broker, observer, and audit dependency failures; and executed-path/post-effect audit correlation. The 57-test focused suite includes dedicated negative regressions for these classes. These corrections are meaningful CE-1 implementation evidence, not exhaustive security assurance.

Security and evidence boundaries

- The broker capability, private target method, and exact environment type are

application-level Python controls. They are not OS/process isolation, a reference monitor, or protection against arbitrary hostile same-process code.

- Request and authorization ledgers are process-local memory. They are not

durable, distributed, crash consistent, or sufficient for multi-node replay prevention/idempotency.

- Evidence attestations use domain-separated runtime HMAC synthetic keys. They do not establish

enterprise device identity, PKI/HSM custody, source authenticity, nonrepudiation, rotation/revocation, or independent provenance.

- Decision and target verifiers are functionally separate non-model components

in the same project/process. They are not an external oracle or organizationally independent assurance.

- Human approval authorizes a reevaluation receipt for an exact scope only. It

is not an action token, cannot execute, and does not itself cause reevaluation.

- Audit is self-custodied and hash linked, not externally anchored, WORM

protected, trusted-time bound, or resistant to complete authorized replacement/truncation.

- Two synthetic targets and injected faults do not establish vendor semantics,

real topology, eventual consistency, rollback feasibility, human workflow, mission outcome, or operational safety.

Reproduce the local Phase 3 observations

```
PYTHONDONTWRITEBYTECODE=1 PYTHONPATH=src python3 -m unittest \
tests.test_phase3_contracts \
tests.test_phase3_decision_path \
tests.test_phase3_authorization_boundary \
tests.test_phase3_adversarial \
tests.test_phase3_end_to_end \
tests.test_phase3_corpus \
tests.test_phase3_release_blockers -v

demo_dir="$(mktemp -d /tmp/adf-phase3-demo.XXXXXXX)"
PYTHONDONTWRITEBYTECODE=1 PYTHONPATH=src python3 run_phase3.py \
--output-dir "$demo_dir"

corpus_dir="$(mktemp -d /tmp/adf-phase3-corpus.XXXXXXX)"
PYTHONDONTWRITEBYTECODE=1 PYTHONPATH=src python3 run_phase3_corpus.py \
--output-dir "$corpus_dir"
```

Reproduce the settled local repository aggregate:

```
PYTHONDONTWRITEBYTECODE=1 PYTHONPATH=src python3 -m unittest discover -s tests -v
```

Successful local return is diagnostic evidence for the checkout that ran. It does not substitute for an exact commit, CI, manifest verification, or external review.

Forward plan and gates

1. **Close the local candidate.** Reconcile traceability and diagrams and resolve every release-blocking finding.
2. **Freeze and verify.** Run the focused Phase 3 modules, final repository aggregate, both demos, corpus, documentation/link/schema/CSV checks, and integrity/package checks against settled bytes.
3. **Commit and publish deliberately.** Create one exact Phase 3 commit, record its results, publish only under explicit authorization, and require green exact-commit GitHub CI and Dependency Graph checks.
4. **Keep data-bearing work separate.** P2-CE-005, if pursued, retains its two-commit protocol. Historical replay requires an authenticated restricted Gate B package. Neither is implied by Phase 3 simulation success.
5. **Require a new architecture before integration.** A read-only service or controlled non-production action phase requires enterprise identity and source trust, process isolation, managed keys, durable distributed idempotency, vendor-specific broker/readback, rollback/reconciliation, external audit custody, change/incident authority, and authorizing-official acceptance.

See the [Phase 3 index](#), [architecture](#), [security and safety case](#), [T&E plan](#), and [requirements traceability](#).